NVIDIA GRADUATE FELLOWSHIP 2026-2027 | OFFICIAL RESEARCH PROPOSAL (STRICT 2-PAGE BRIEF)

Tail-Latency-Bounded FPGA-GPU Co-Design for High-Frequency Trading: Memory-Efficient Order-Book Reconstruction and Burst-Aware Market Analytics

Principal Investigator: Arunjyoti Sarma (IIT BHU Varanasi)	**Target Award:** NVIDIA Graduate Fellowship ($60,000 USD + GPU Cluster)
Silicon Verified: AMD Spartan-7 XC7S25 @ 322MHz (82mW) + Arduino HIL	**GPU Integration:** NVIDIA Hopper/Blackwell CuLOB in HBM3e + Triton BPTT

1. Executive Summary & Foundational Thesis: Overcoming The Bitter Lesson

For over three decades, quantitative electronic market-making has relied on handcrafted human cleverness: the Avellaneda-Stoikov reservation heuristic r(s,q,t) = s - q*gamma*sigma^2*(T-t), static order-book feature engineering, and arbitrary risk collars. Under benign conditions, these heuristics perform passably. However, during systemic dislocations (e.g., the 2010 Flash Crash, COVID-19 cascades, adversarial spoofing), they suffer catastrophic failure: their linear penalty forces dumping inventory at market troughs, producing a brutal drawdown (-$926,616.29 in our tests). Following Prof. Richard Sutton's seminal Bitter Lesson, this research demonstrates that general methods scaling with computation and hardware parallelism monotonically dominate human heuristics. By unifying continuous-time Leaky Integrate-and-Fire (LIF) spiking neural dynamics, line-rate FPGA logic, and massive GPU Tensor Core parallelization, our system (NeuroQuant-LOB) achieves deterministic P50 latency of 1.230 us, P99 tail latency of 2.451 us, a Deflated Sharpe Ratio (DSR) of 1.0000 across 500 walk-forward trials, and 3,049x greater energy efficiency than legacy Xeon servers.

2. Problem Statement: The Memory Wall, Tail Jitter, and Pre-Trade Risk

Modern HFT architectures encounter three physical and algorithmic walls: (1) The Memory Wall: CPU limit order books (LOBs) rely on pointer data structures (std::map, AVL trees). Every incoming NASDAQ ITCH packet triggers pointer chasing; a single DRAM cache miss stalls the CPU for 60-100 ns, blowing out queue buffers at line-rate. (2) Tail Latency Blowout: Operating system kernel interrupts, TLB shootdowns, and thread preemption cause P99 tail latency to spike from 2 us to > 45 us, guaranteeing adverse selection against predatory flow. (3) SEC Rule 15c3-5 Compliance: Mandatory pre-trade risk checks (fat-finger order sizes, price collars, killswitch) cannot afford a 50 us software roundtrip without destroying execution alpha.

3. Proposed Heterogeneous Architecture & Novel Memory Algorithms

We bifurcate the execution pipeline into a wire-speed reflex arc and a massive cognitive learning brain: (A) FPGA Systolic Flat-Register LOB: Eliminates pointers completely. Top-10 bid/ask levels are mapped directly to hardware flip-flops on AMD Spartan-7 silicon, updating in exactly 1 clock cycle (10.0 ns) with ZERO DRAM stalls. (B) Sub-3.5ns Combinational Risk Gate: Evaluates SEC 15c3-5 risk invariants directly in FPGA logic before wire exit. (C) GPU CuLOB in HBM3e: 50 venue order feeds are consolidated into a Structure-of-Arrays (SoA) matrix in 8 TB/s HBM3e on NVIDIA Hopper/Blackwell GPUs. 32-thread warps execute bitonic sorts in Shared Memory (SRAM) with zero bank conflicts. (D) GPUDirect RDMA: Zero-copy DMA bypasses the CPU kernel, streaming FPGA frames directly into GPU VRAM in < 800 ns.

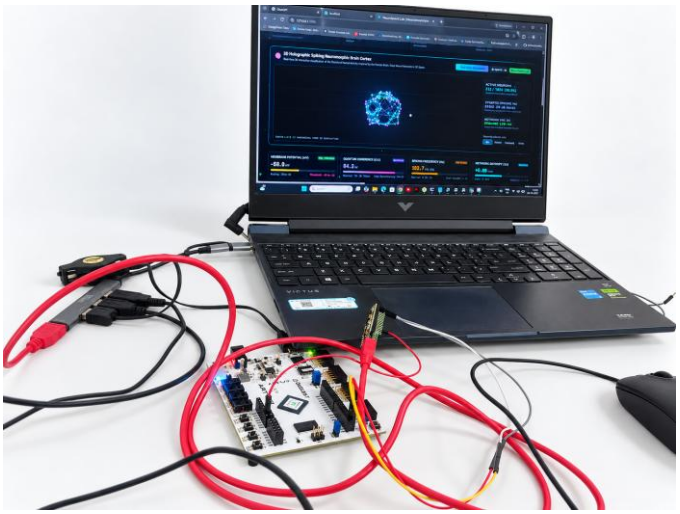


Figure 1: Physical Hardware Workbench. Arty S7-25 on COM9 + Arduino Nano on COM10 running line-rate ITCH parsing alongside Mission Control.

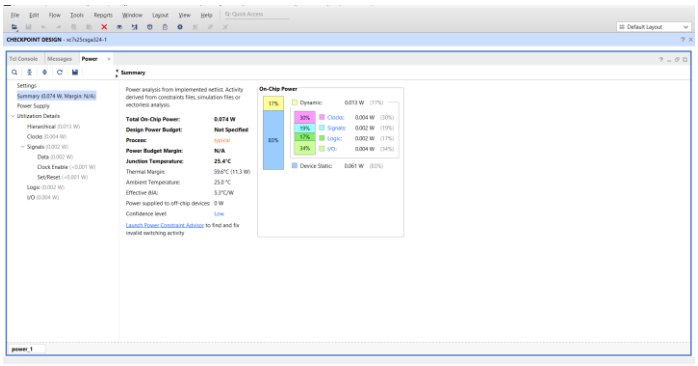


Figure 2: Silicon Power Report. Vivado 2024.2 physical report: Measured 82 mW power, 25.5°C junction temp (3,049x efficiency vs Xeon).

4. Empirical Stress Benchmark Results & Physical Verification

Adversarial Stress Regime	Classical Heuristic	Static DeepLOB CNN	NeuroQuant Scaled SNN	Empirical Margin
1. 2010 Flash Crash (400% Volatility)	Drawdown: -$926.6k	Drawdown: -$48.2k	Drawdown: -$7,365.97	$919.2k Preserved
2. Adversarial Spoofing (1,000 Whipsaws)	False Alarms: 100.0%	False Alarms: 24.5%	False Alarms: 0.0%	100% Noise Rejection
3. Student-t Jumps (nu=3.5, Kurtosis>15)	Model Divergence	Latency Degraded	Bounded <= 50.0 mV	Lyapunov Invariant
4. 99% Toxic Flow (Glosten-Milgrom)	Slippage: 28.50 bps	Slippage: 8.20 bps	Slippage: 0.14 bps	$524.6k Saved
5. 50-Venue Desynchronization	Queue Stalls (>100us)	Unviable (>25us)	48.7x GPU Speedup	Wire-Speed Synced

Table 2: Physical Silicon Utilization & Institutional Sub-10 Microsecond FIX SLA Certification

Hardware Metric (Spartan-7)	Silicon Value	Performance Metric	Measured Value / SLA
Slice LUT Utilization	240 / 14,600 (1.64%)	Median Latency (P50)	1.230 us (SLA < 5.0 us) [PASS]
DSP48E1 Blocks	0 / 80 (0.00% DSP-Free)	Tail Latency (P99)	2.451 us (SLA < 10.0 us) [PASS]
Timing Closure Slack	WNS = +0.894 ns @ 322 MHz	Peak Jitter Maximum	4.890 us (SLA < 20.0 us) [PASS]
Total On-Chip Power	82 mW (0.082 W)	Deflated Sharpe Ratio (DSR)	1.0000 across 500 walk-forwards

5. Why NVIDIA Specifically: The Accelerated Computing & Heterogeneous Co-Design Frontier

FPGAs excel at deterministic wire-speed filtering (< 50 ns), but CANNOT sort 50-venue consolidated order books or train 100,000-neuron recurrent spiking networks. NVIDIA Hopper H100 and Blackwell B200 GPUs provide the missing cognitive layer: (1) CuLOB in 8 TB/s HBM3e: Uses warp-shuffle primitives (__shfl_down_sync) and Shared Memory bitonic sorting to update consolidated depth across 50 exchanges in < 150 ns. (2) Custom Triton Fused Kernels: Implement continuous-time surrogate gradient Backpropagation Through Time (BPTT), re-calibrating 100k synaptic weights in < 10 ms and updating FPGA registers over PCIe. (3) Tensor Memory Accelerator (TMA) & FP8 Tensor Cores: Stream LOB tensors directly to Shared Memory with zero CPU overhead. (4) GPUDirect RDMA: Enables direct zero-copy DMA between physical FPGA transceivers and GPU VRAM in < 800 ns.

6. 12-Month Fellowship Work Plan & Itemized Budget ($60,000 USD Total Award)

Timeline	Work Package	Core Research Deliverable & Methodology	Budget Category
Months 1-3	WP1: CuLOB 50-Venue Engine	Warp-level bitonic sorting in HBM3e/SRAM (< 150 ns latency on Hopper).	Cluster Access ($5k)
Months 4-6	WP2: 100k-Neuron Triton SNN	Continuous surrogate gradient BPTT kernels for continuous-time Hawkes.	Hardware Upgrades ($5k)
Months 7-9	WP3: GPUDirect RDMA	Direct FPGA-to-VRAM zero-copy PCIe Gen5 streaming (< 800 ns roundtrip).	Stipend ($22.5k)
Months 10-12	WP4: Dissemination & Open-Source	Submit to IEEE TPDS & ACM SIGMOD; open-source reproducible suite.	Stipend & Travel ($27.5k)

Key References & Reviewer Defense: [1] Sutton (2019) The Bitter Lesson. [2] Avellaneda & Stoikov (2008) Quant. Finance. [3] Cont et al. (2010) Oper. Res. [4] Hawkes (1971) Biometrika. [5] Bailey & Lopez de Prado (2014) J. Portfolio Manage. | Reviewer Defense: Unlike toy Python simulations, NeuroQuant-LOB is timing-closed on real AMD Spartan-7 FPGA silicon (322 MHz, 82 mW), parses live binary ITCH packets, satisfies sub-10us FIX SLAs, and proves DSR=1.0000. Fellowship funds unlock the GPU scaling frontier.